



scala-miniboxing.org

21st of October 2014
PDXScala Meetup
Portland, OR

Vlad URECHE

PhD student in the Scala Team @ EPFL

Miniboxing guy. Also worked on Scala specialization, the backend and scaladoc.



@VladUreche




@VladUreche



vlad.ureche@epfl.ch



scala-miniboxing.org



scala-miniboxing.org



Miniboxing

Theory

Practice

Benchmarks

Conclusion




Erased Generics

```
def identity[T](t: T): T = t
```

Erased Generics


```
def identity[T](t: T): T = t
```



scalac / javac

Erased Generics

```
def identity[T](t: T): T = t
```



scalac / javac


```
def identity(t: Object): Object = t
```

Erased Generics

identity(5)

Erased Generics


identity(5)



scalac / javac

Erased Generics

identity(5)




scalac / javac

identity(j.l.Integer.valueOf(5)).intValue

Erased Generics

identity(5)




scalac / javac

identity(j.l.Integer.valueOf(5)).intValue


Object representation

Erased Generics

identity(5)



scalac / javac




identity(j.l.Integer.valueOf(5)).intValue



Object representation

Erased Generics

identity(5)



scalac / javac




identity(j.l.Integer.valueOf(5)).intValue

Object representation

Erased Generics

identity(5)



scalac / javac

inflates heap requirements

produces garbage


identity(j.l.Integer.valueOf(5)).intValue

Object representation

indirect
(slow) access
to the value

Erased Generics

identity(5)



scalac / javac

inflates heap requirements

produces garbage

identity(j.l.Integer.valueOf(5)).intValue

Object representation

indirect
(slow) access
to the value


breaks locality
guarantees

Specialization

```
def identity[T](t: T): T = t
```


Specialization


```
def identity[T](t: T): T = t
```



specialization

Specialization

```
def identity[T](t: T): T = t
```




specialization

```
def identity(t: Object): Object = t
```

Specialization

```
def identity[T](t: T): T = t
```




specialization

```
def identity(t: Object): Object = t  
def identity_Z(t: bool): bool = t
```

Specialization

```
def identity[T](t: T): T = t
```




specialization

```
def identity(t: Object): Object = t  
def identity_Z(t: bool): bool = t  
def identity_C(t: char): char = t
```

Specialization

```
def identity[T](t: T): T = t
```



specialization


```
def identity(t: Object): Object = t  
def identity_Z(t: bool): bool = t  
def identity_C(t: char): char = t  
... (7 other variants)
```

Specialization

identity(5)

Specialization


identity(5)



specialization

Specialization

identity(5)




specialization

identity_1(5)

Specialization

identity(5)




specialization

identity_I(5)

The variant of identity
specialized for **int**

Specialization

identity(5)



specialization

identity_I(5) // no boxing!


The variant of identity
specialized for **int**

Specialization

```
def tupled[T1, T2](t1: T1, t2: T2) ...
```

Specialization


```
def tupled[T1, T2](t1: T1, t2: T2) ...
```



specialization

Specialization

```
def tupled[T1, T2](t1: T1, t2: T2) ...
```




specialization

```
// 100 methods ( $10^2$ )
```


Specialization

```
def tupled[T1, T2](t1: T1, t2: T2) ...
```




specialization

```
// 100 methods ( $10^2$ )
```



Specialization

```
def tupled[T1, T2](t1: T1, t2: T2) ...
```




specialization

```
// 100 methods ( $10^2$ )
```




Can we do
better?

Miniboxing




```
def identity[T](t: T): T = t
```


Miniboxing



```
def identity[T](t: T): T = t
```




miniboxing

Miniboxing



```
def identity[T](t: T): T = t
```




miniboxing

```
def identity(t: Object): Object = t
```

Miniboxing




```
def identity[T](t: T): T = t
```




miniboxing

```
def identity(t: Object): Object = t  
def identity_M(..., t: long): long = t
```

Miniboxing



```
def identity[T](t: T): T = t
```




miniboxing


```
def identity(t: Object): Object = t  
def identity_M(..., t: long): long = t
```

long **encodes** all
primitive types

Miniboxing



```
def identity[T](t: T): T = t
```



miniboxing

```
def identity(t: Object): Object = t  
def identity_M(..., t: long): long = t
```

long **encodes** all
primitive types




Only 2^n variants

Miniboxing




identity(3)

Miniboxing




identity(3)




miniboxing

scala-miniboxing.org

Miniboxing




identity(3)




miniboxing

identity_M(..., int2minibox(3))

Miniboxing



identity(3)




miniboxing


identity_M(..., int2minibox(3))

The miniboxed
variant of identity

Miniboxing



identity(3)



miniboxing

identity_M(..., int2minibox(3))

The miniboxed
variant of identity

Encoding the integer
into a long integer



Miniboxing

Theory

Practice

Benchmarks

Conclusion



scala-miniboxing.org

○ The Theory of Miniboxing

○ Class Transformation

○ Late Data Layout

○ ...




○ The Theory of Miniboxing

● Class Transformation

○ Late Data Layout

○ ...




class C[@specialized T](t: T)


class C[@specialized T](t: T)



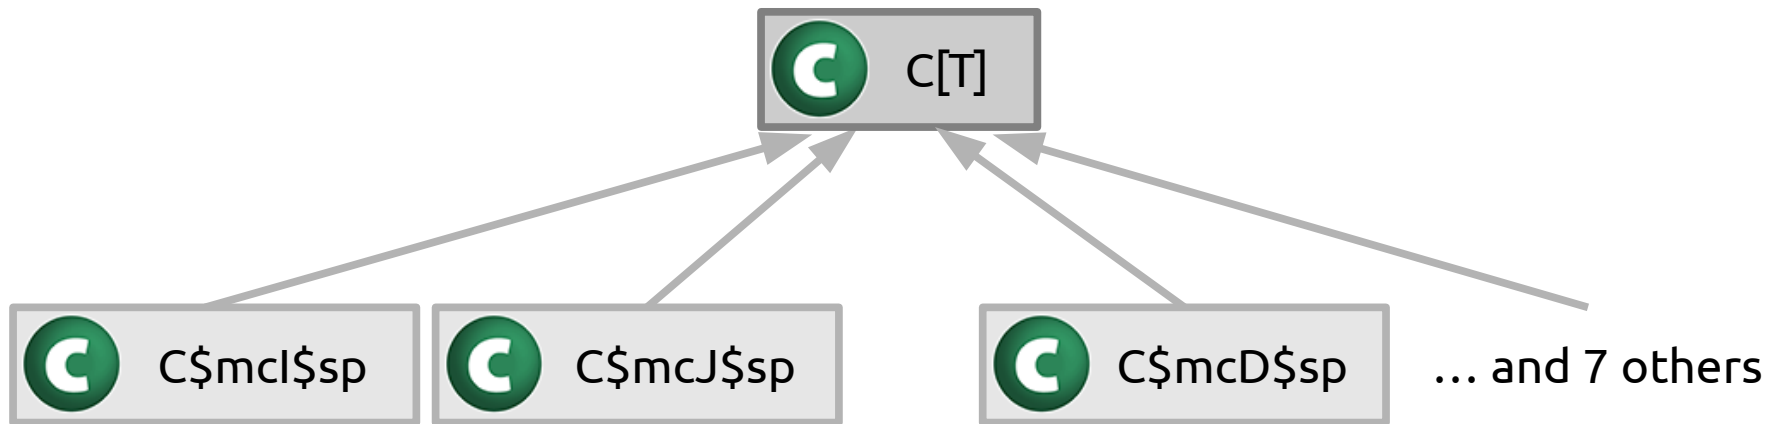
class C[@specialized T](t: T)



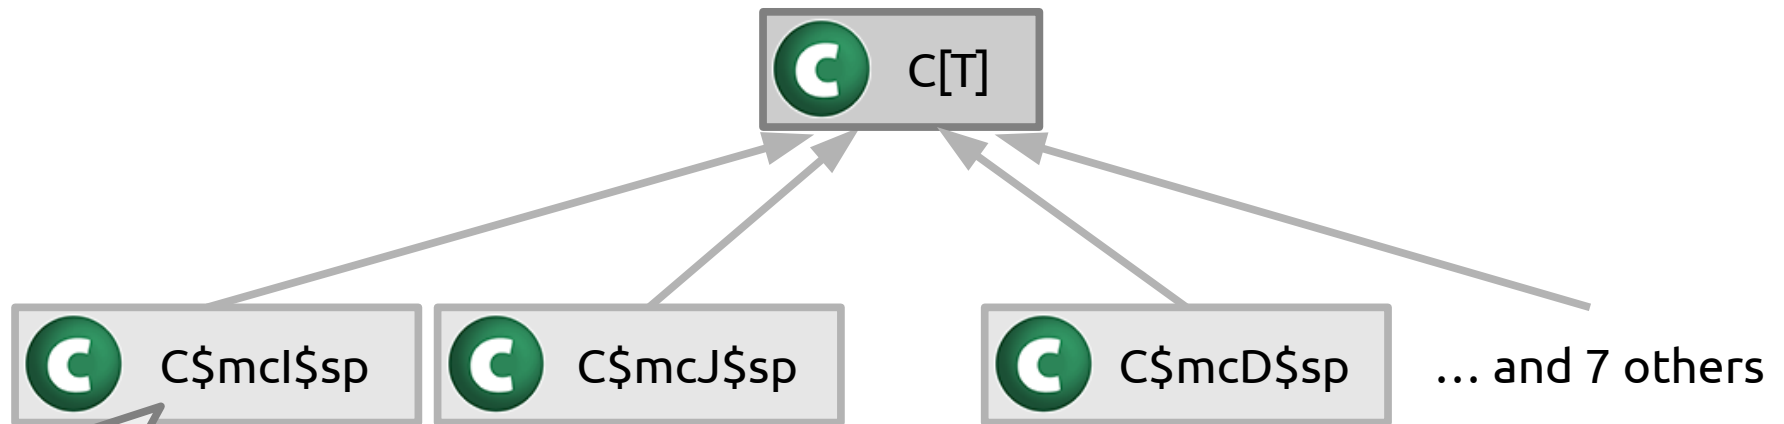
class C[@specialized T](t: T)



class C[@specialized T](t: T)

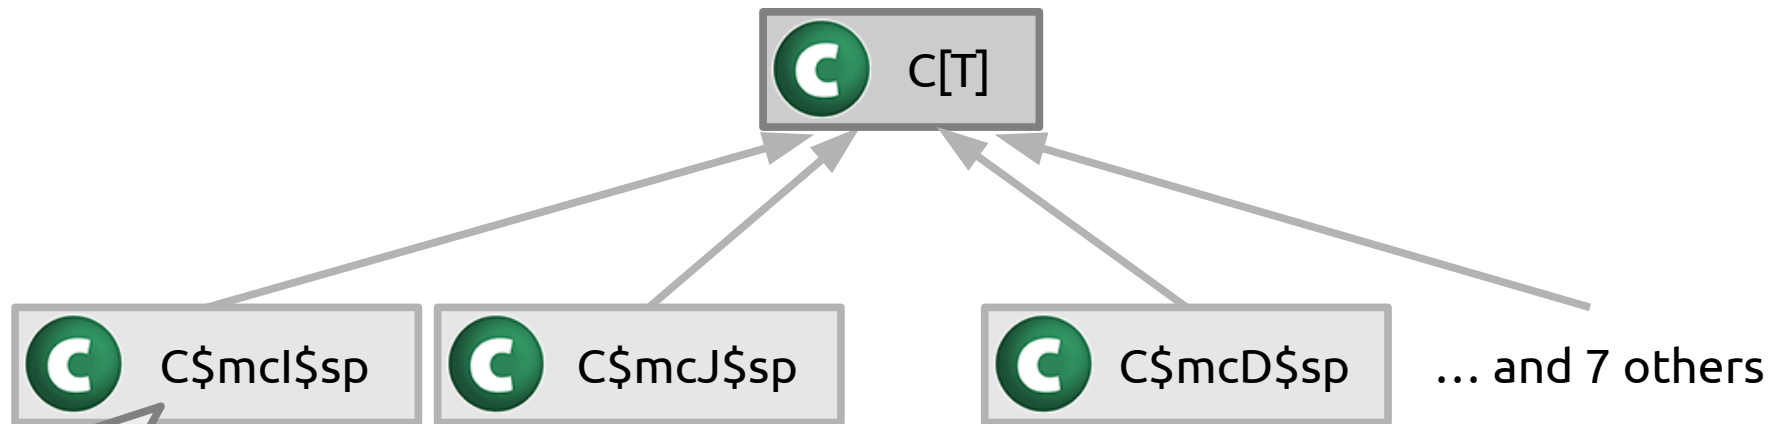


class C[@specialized T](t: T)



What is the field t
In this class?

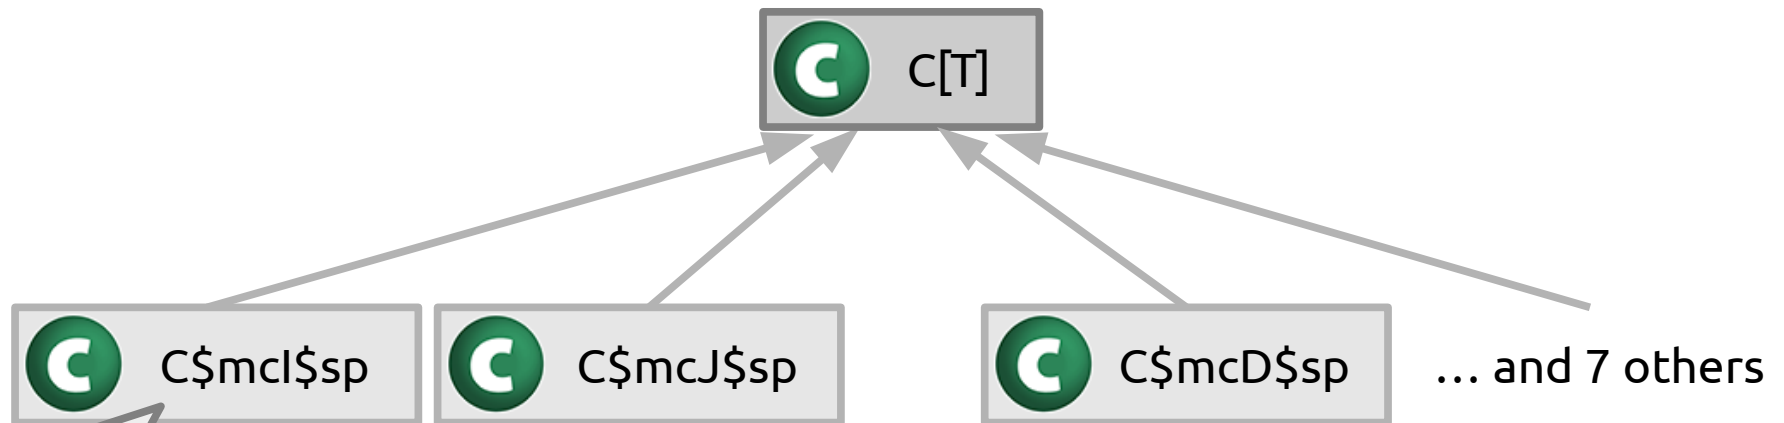
class C[@specialized T](t: T)



What is the field `t`
In this class?

Actually there are two
fields: a generic one from
`C` and a specialized one

class C[@specialized T](t: T)



What is the field `t`
In this class?

Actually there are two
fields: a generic one from
C and a specialized one

That's bad:
we waste memory

SI-3585

class C[@miniboxed T](t: T)


class C[@miniboxed T](t: T)




class C[@miniboxed T](t: T)




class C[@miniboxed T](t: T)



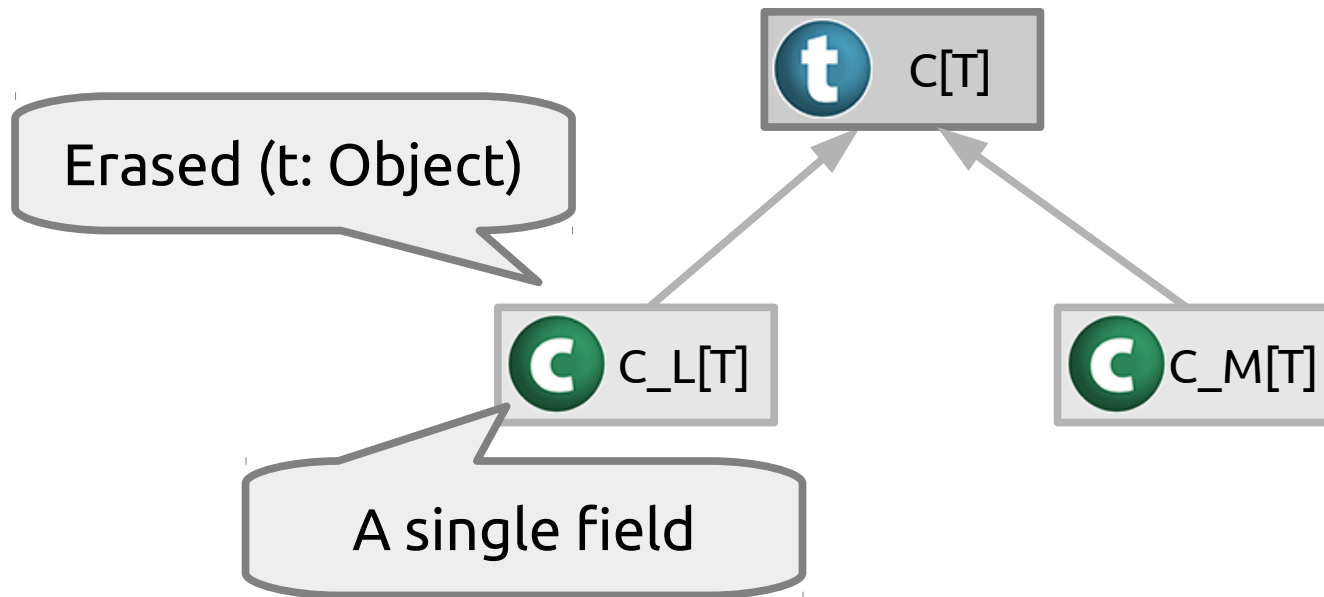
class C[@miniboxed T](t: T)




class C[@miniboxed T](t: T)



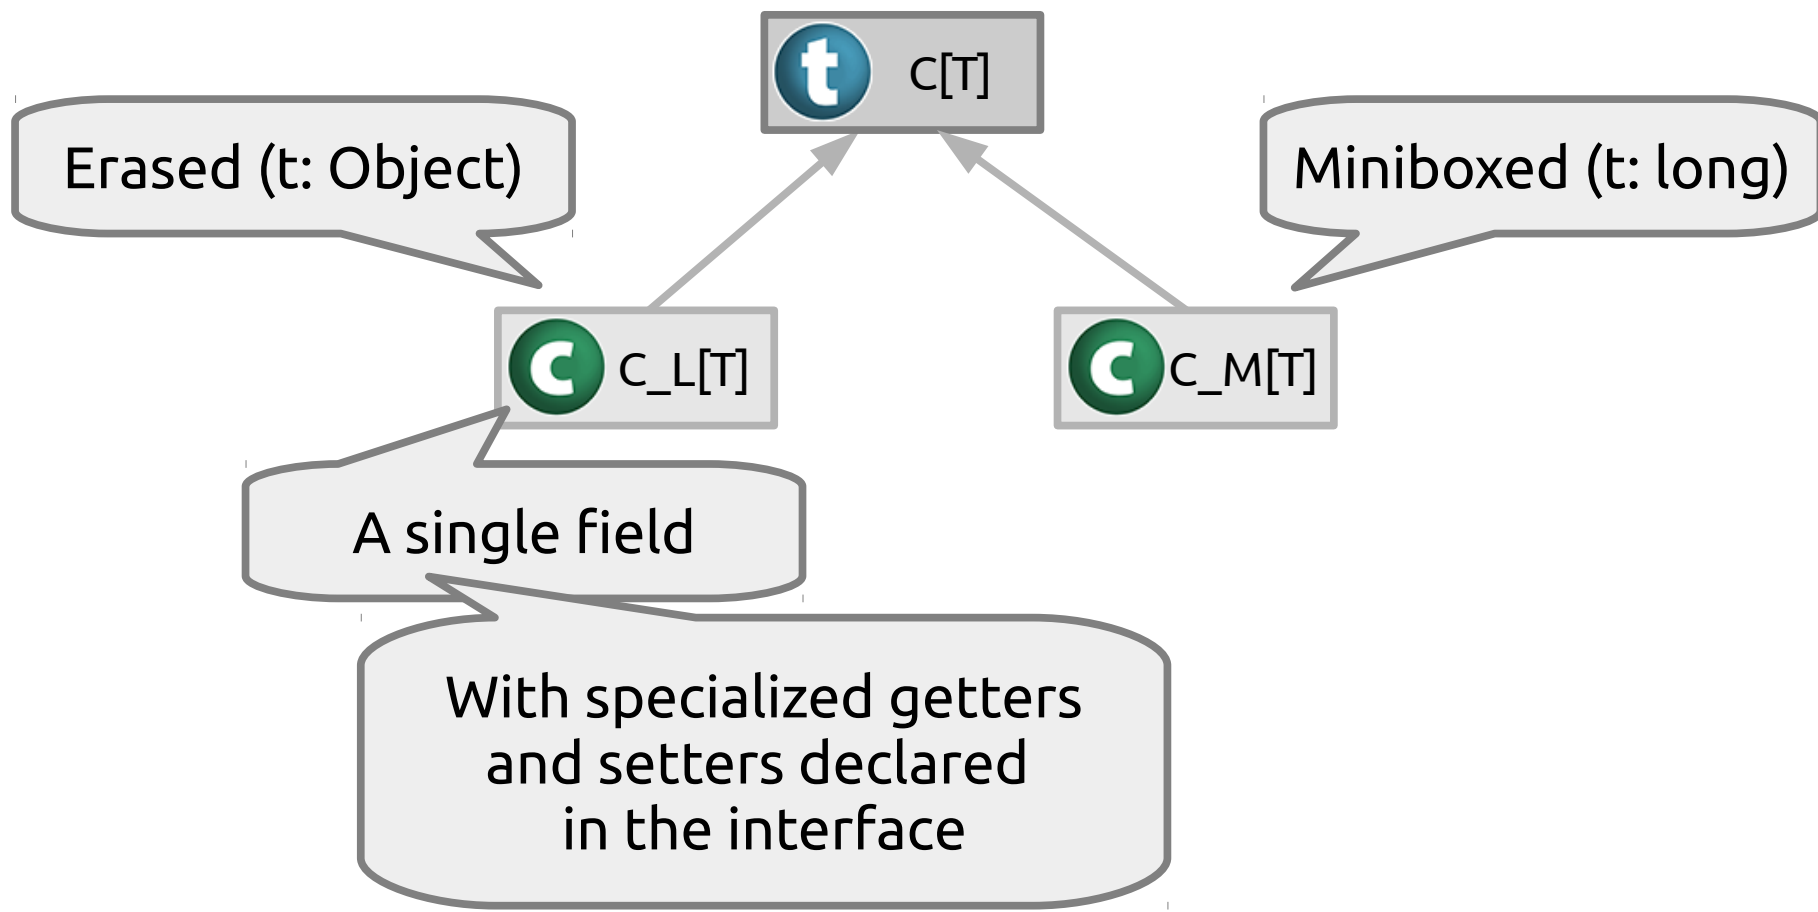
class C[@miniboxed T](t: T)



class C[@miniboxed T](t: T)



class C[@miniboxed T](t: T)




```
class D[@specialized T](t: T)  
    extends C[T]
```

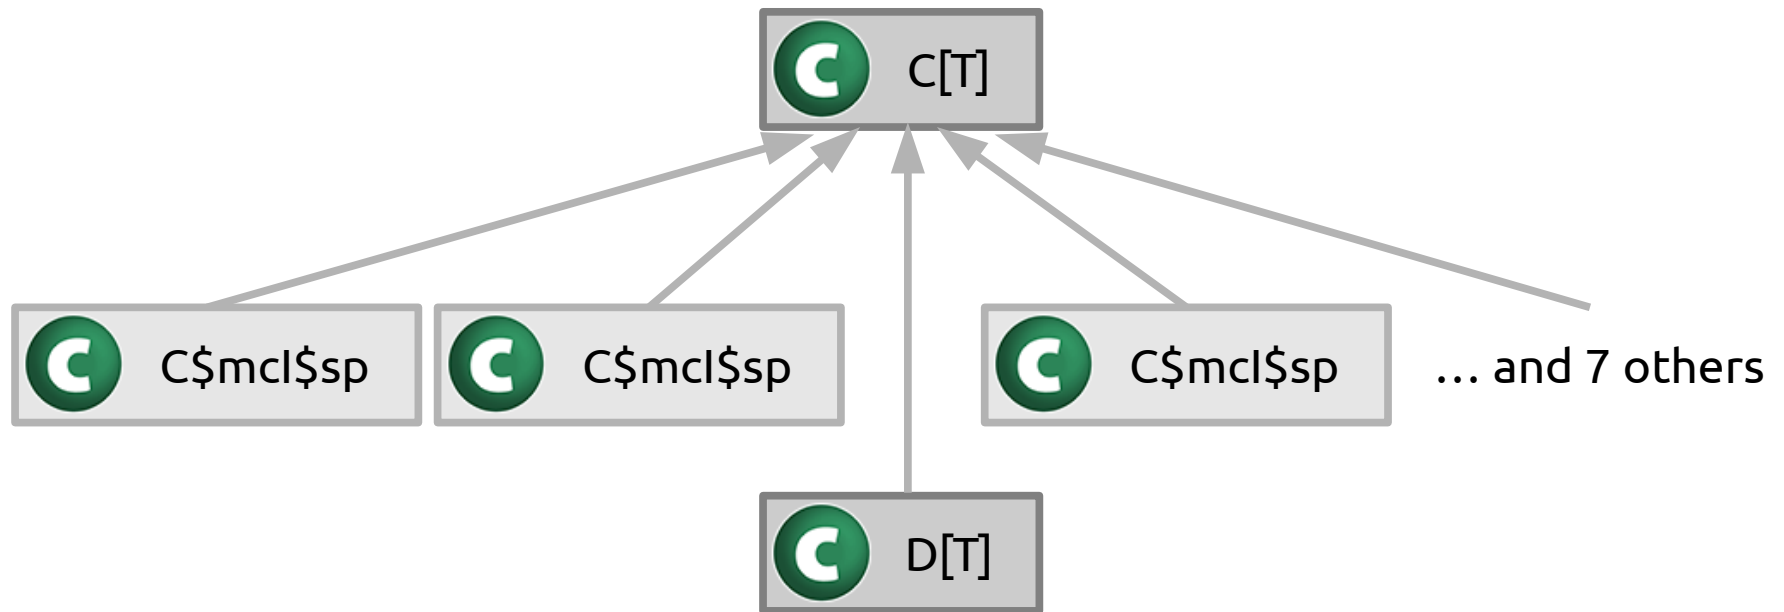
```
class D[@specialized T](t: T)  
  extends C[T]
```



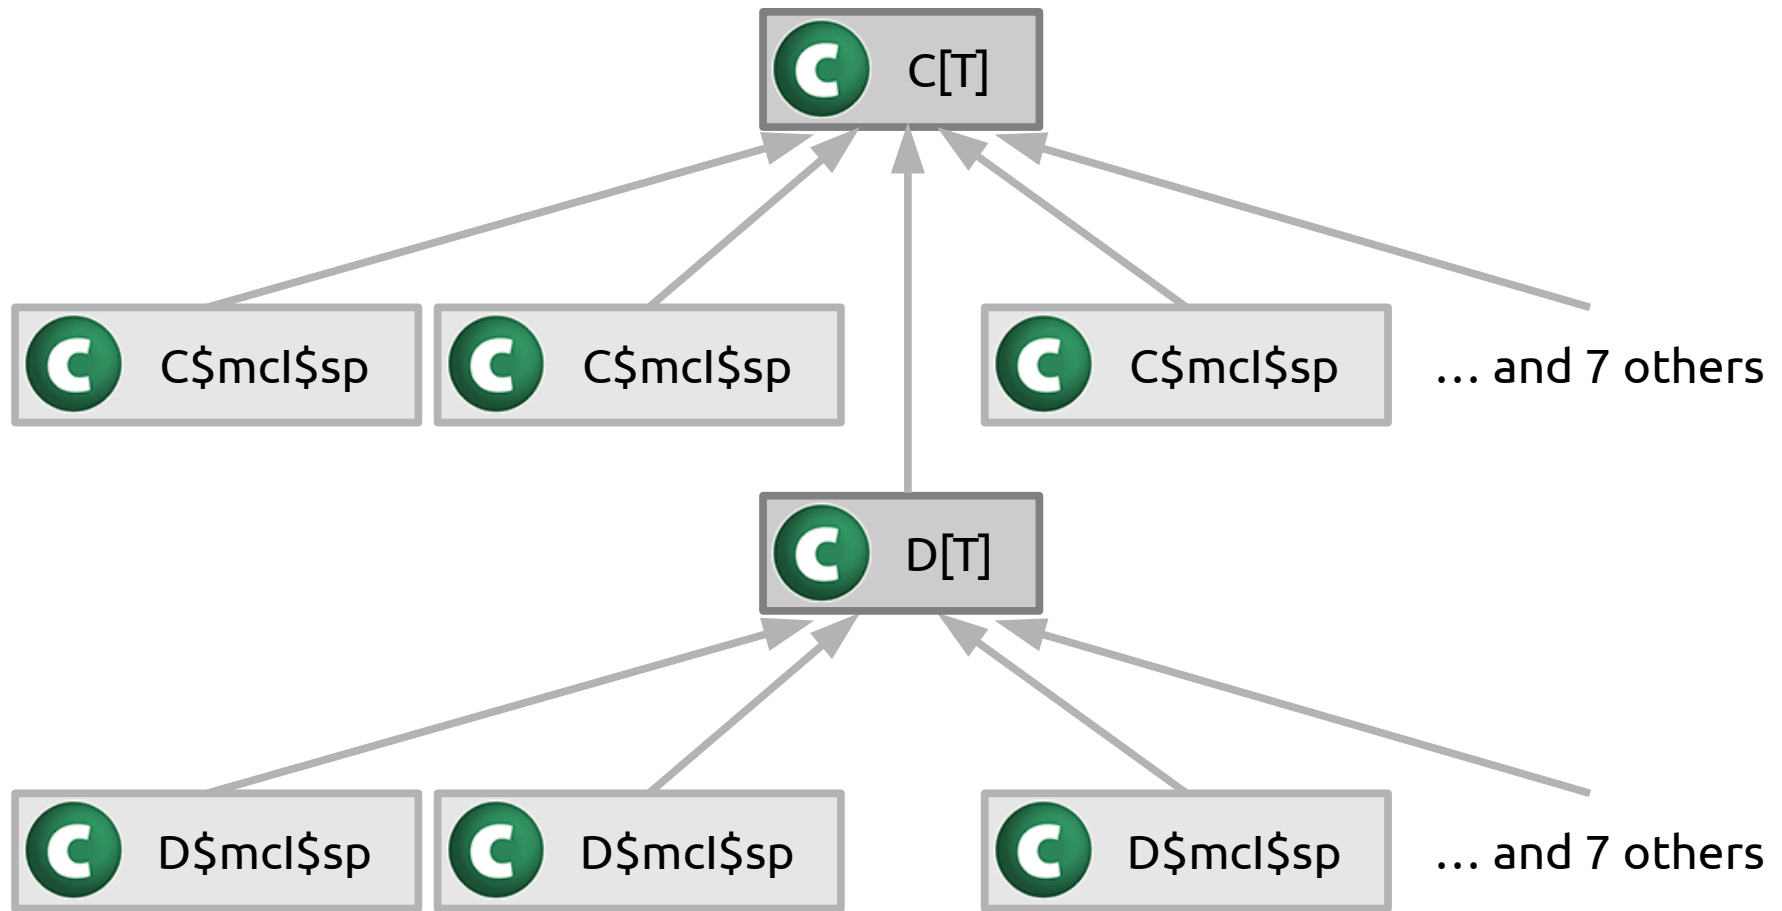

```
class D[@specialized T](t: T)  
  extends C[T]
```



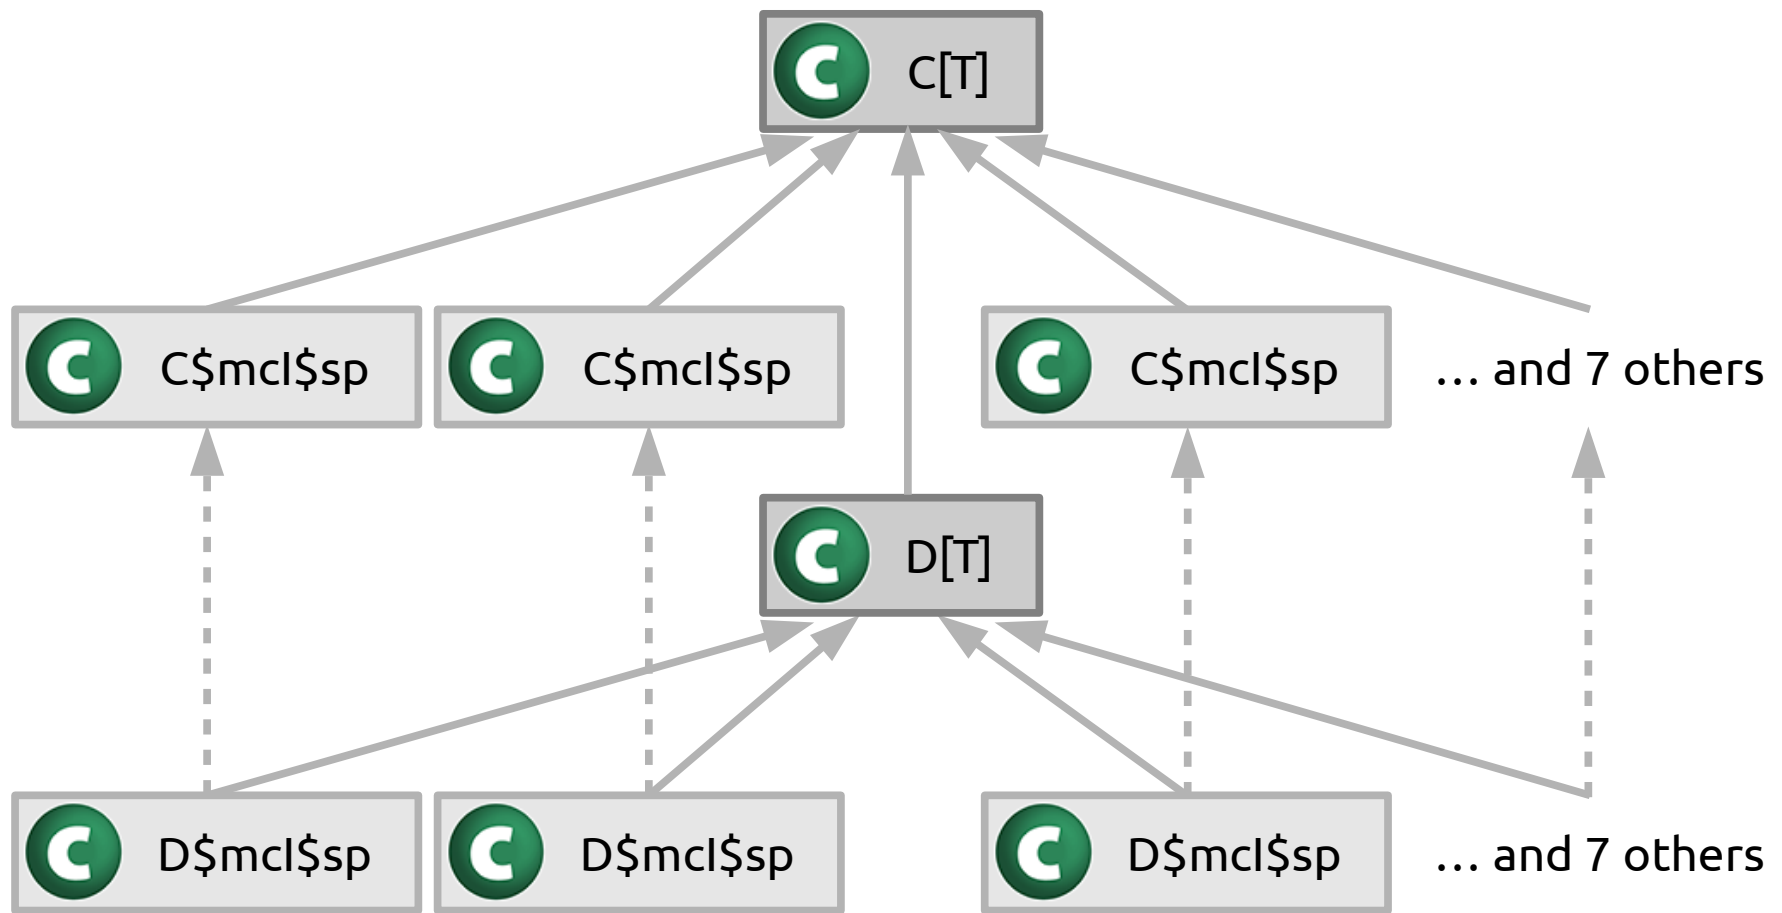
**class D[@specialized T](t: T)
extends C[T]**



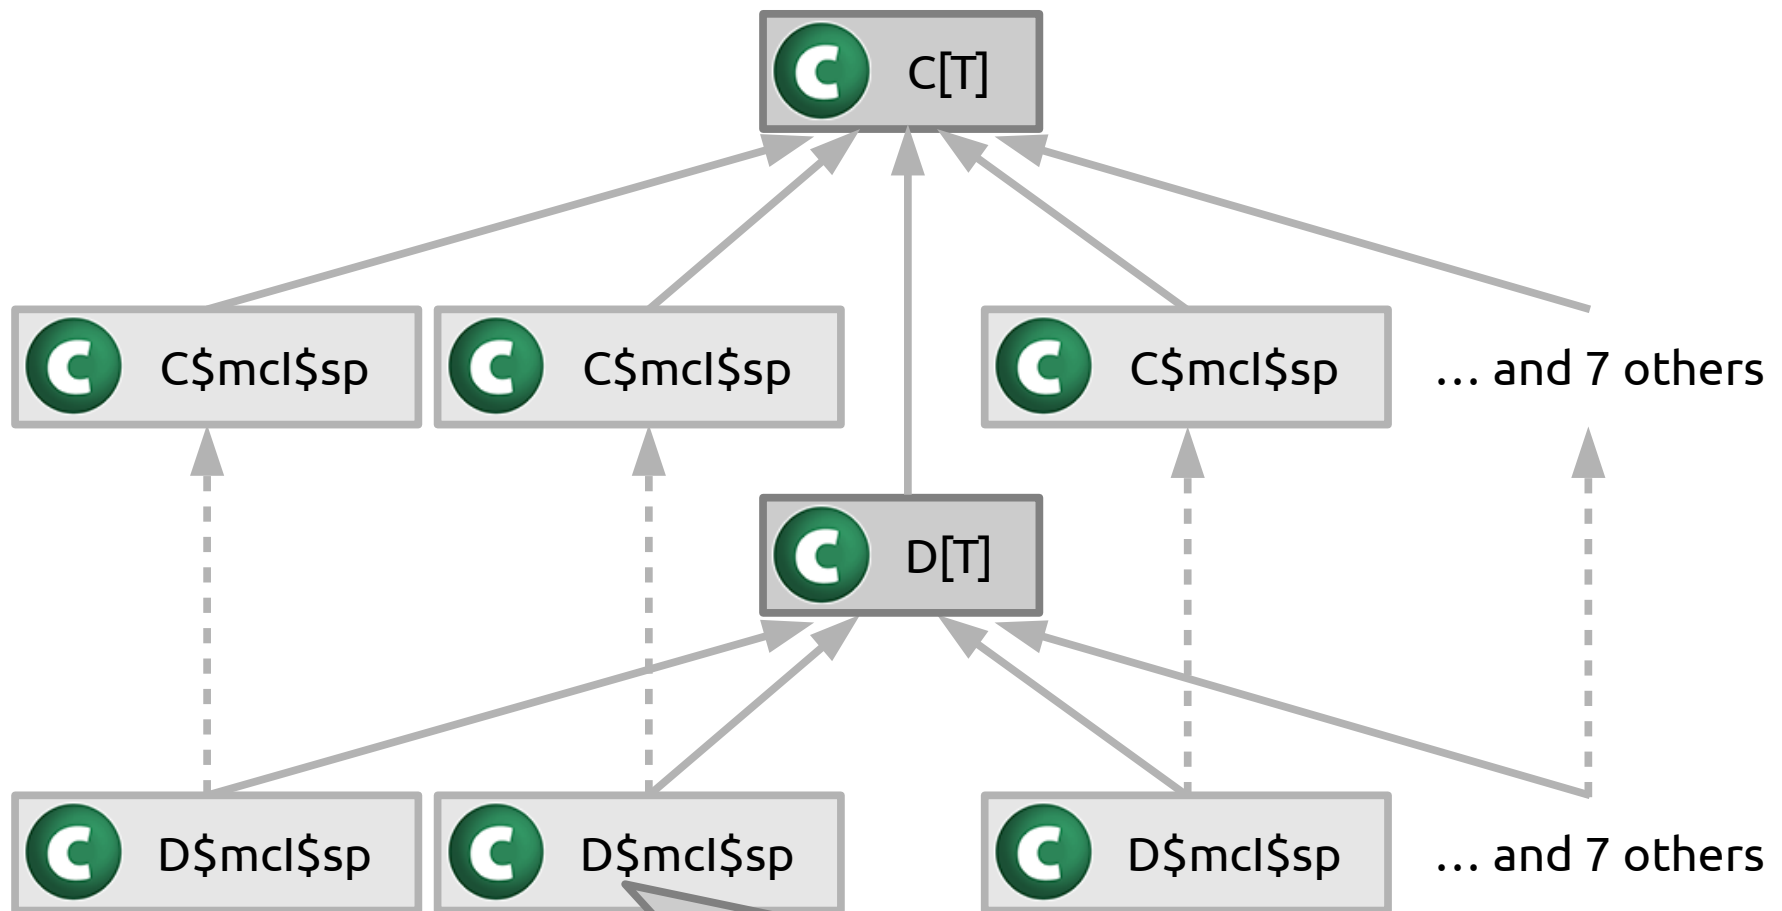
```
class D[@specialized T](t: T)  
  extends C[T]
```



**class D[@specialized T](t: T)
extends C[T]**



class D[@specialized T](t: T) extends C[T]

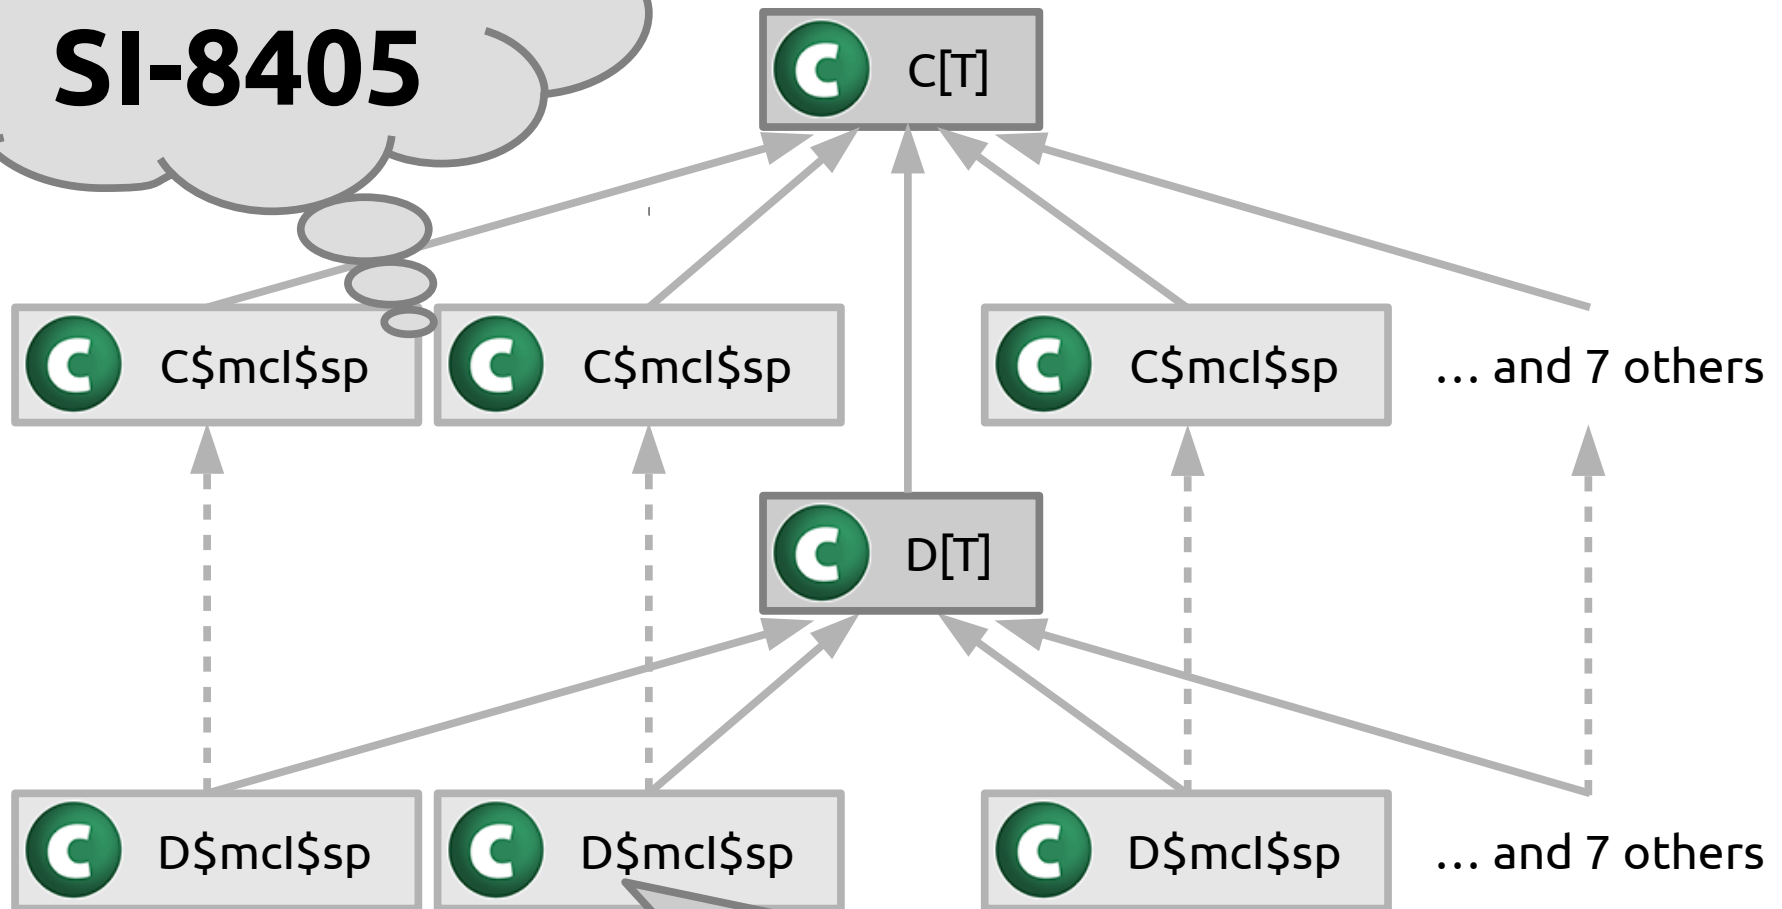


NoSuchThingError: Multiple Inheritance on the JVM

class D[@specialized T](t: T) extends C[T]

Class inheritance
doesn't work


SI-8405




NoSuchThingError: Multiple Inheritance on the JVM

```
class D[@miniboxed T](t: T)  
    extends C[T]
```

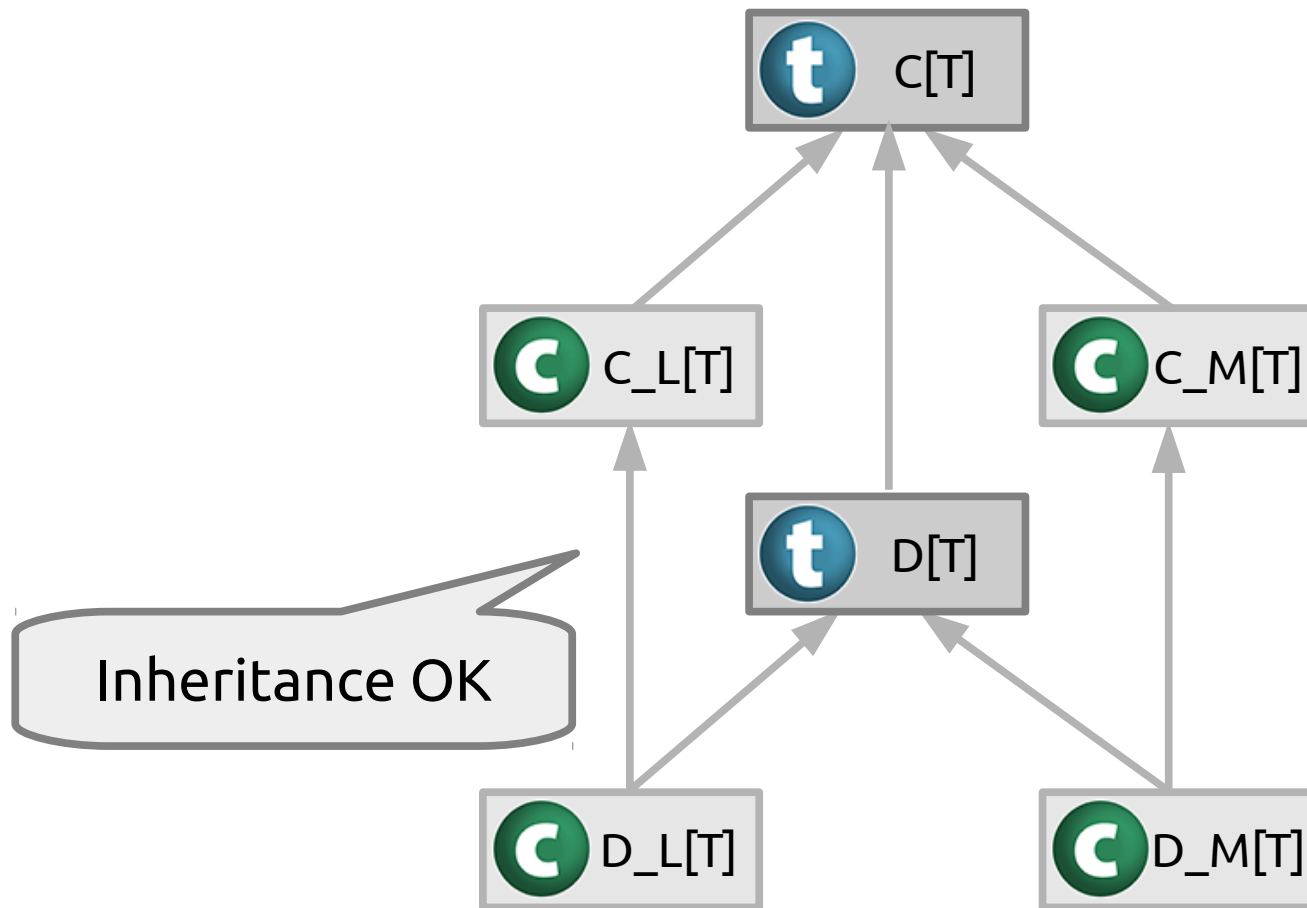
```
class D[@miniboxed T](t: T)  
  extends C[T]
```




```
class D[@miniboxed T](t: T)  
  extends C[T]
```



```
class D[@miniboxed T](t: T)  
  extends C[T]
```



○ The Theory of Miniboxing

○ Class Transformation

● Late Data Layout

○ ...



Late Data Layout

Late Data Layout: Unifying Data Representation Transformations

Vlad Ureche Eugene Burmako Martin Odersky
EPFL, Switzerland
{first.last}@epfl.ch



Abstract

Values need to be represented differently when interacting with certain language features. For example, an integer has to take an object-based representation when interacting with erased generics, although, for performance reasons, the stack-based value representation is better. To abstract over these implementation details, some programming languages choose to expose a unified high-level concept (the integer) and let the compiler choose its exact representation and insert coercions where necessary.

This pattern appears in multiple language features such as value classes, specialization and multi-stage programming: they all expose a unified concept which they later refine into multiple representations. Yet, the underlying compiler implementations typically entangle the core mechanism with assumptions about the alternative representations and their interaction with other language features.

In this paper we present the Late Data Layout mechanism, a simple but versatile type-driven generalization that subsumes and improves the state-of-the-art representation transformations. In doing so, we make two key observations: (1) annotated types conveniently capture the semantics of using multiple representations and (2) local type inference can be used to consistently and optimally introduce coercions.

We validated our approach by implementing three language features as Scala compiler extensions: value classes, specialization (using the miniboxing representation) and a simplified multi-stage programming mechanism.

Categories and Subject Descriptors: D.3.3 [Language Constructs and Features]: Polymorphism; E.2 [Object representation]

Keywords: Data Representation; Object-oriented; Annotated Types; Type Systems; Local Type Inference

1. Introduction

Language and compiler designers are well aware of the intricacies of erased generics [15, 21, 30, 32, 35, 42, 46, 75], one of which is requiring object-based representations for primitive types. To illustrate this, let us analyze the `identity` method, parameterized on the argument type, `T`:

```
def identity[T](arg: T): T = arg
val x: Int = identity[Int](5)
```

The low-level compiled code for `identity` needs to handle incoming arguments of different sizes and semantics: booleans, bytes, characters, integers, floating point numbers and references to heap-allocated objects. To implement this, some compilers impose a uniform representation, usually based on references to heap objects. This means that primitive types, such as integers, have to be represented as objects when passed to generic methods. The process of representing primitive types as objects is called boxing. Since boxing slows down execution, whenever primitive types are used outside generic environments, they use their stack-based (unboxed) representation. Thus, in the low-level compiled code, `x` is using the unboxed representation, denoted as `Int`:

```
def identity(arg: Object): Object = arg
// val x: Int = identity[Int](5)
val arg_boxed: Object = box(5)
val ret_boxed: Object = identity(arg_boxed)
val x: Int = unbox(ret_boxed)
```


The low-level code shows the two representations of the high-level `Int` concept: the unboxed primitive `Int` and the boxed `Object`, which is compatible with erased generics. There are two approaches to exposing this duality in programming languages: In Java, both representations are accessible to programmers, making them responsible for the choice and exposing the language feature interactions. On the other hand, in order to avoid burdening programmers with implementation details, languages such as ML, Haskell and Scala expose a unified concept, regardless of its representation. Then, during compilation, the representation is automatically chosen based on the interaction with the other language features and the necessary coercions between representations, such as `box` and `unbox`, are added to the code.

This strategy of exposing a unified high-level concept with multiple representations is used in other language features as well:

- theory of the transformation
 - how **T** becomes **long**
- generalizes miniboxing
 - value classes
 - staging support
 - function representation
 - etc
- **scala-ldl.org**

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, to republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

ACM 174, October 19–21, 2014, Portland, OR, USA.
Copyright 2014 by the owner(s). Publication rights licensed to ACM.
978-1-4503-2581-0/14/01...\$15.00.
http://dx.doi.org/10.1145/2660193.2660197




○ The Theory of Miniboxing

○ Class Transformation

○ Late Data Layout

○ ...






Miniboxing

Theory

Practice

Benchmarks

Conclusion




scala-miniboxing.org

○ Miniboxing in Practice

○ Documentation

○ Quirks

○ Live Demo



○ Miniboxing in Practice

● Documentation

○ Quirks

○ Live Demo





Start here

[Start Page](#)
[Introduction](#)
[Full Tutorial](#)
[Benchmarks](#)

Use it

[Sbt config](#)
[Command line](#)
[Source code](#)
[Issue tracker](#)
[License](#)

Experiment

[Sbt project](#)
[Linked List](#)
[Functions](#)
[Reverse](#)

Discuss

[Mailing List](#)
[Twitter](#) 
[News](#) 

Introduction

Miniboxing is a compilation scheme that improves the performance of generics in [the Scala programming language](#). The miniboxing transformation is generic enough to be potentially useful for any statically typed language running on one of the Java Virtual Machines, such as [Managed X10](#), [Kotlin](#) or [Ceylon](#).

We'll start by following what happens to a generic class, as it gets compiled. Let's take `class C` as an example:

```
class C[T](t: T)
```

After compiling this class to Java Virtual Machine bytecode, under the [erasure transformation](#) one would get bytecode which roughly corresponds to:

```
class C {  
  var t: Object = _      // field  
  def C(t: Object): C = { // constructor  
    this.t = t  
  }  
}
```

As you can see, erasure transformed `t` from a generic value into a pointer to a heap object. While this is perfectly suited for storing a string or another class inside `class C`, it becomes suboptimal when dealing with primitive value types, such as booleans, bytes, integers and floating point numbers.

The reason it's suboptimal is because primitive value types are not heap objects but values passed on the stack, so under the Java Virtual Machine it is common to encode them as heap objects, a process known as [boxing](#).



Start here

[Start Page](#)
[Introduction](#)
[Full Tutorial](#)
[Benchmarks](#)

Use it

[Sbt config](#)
[Command line](#)
[Source code](#)
[Issue tracker](#)
[License](#)

Experiment

[Sbt project](#)
[Linked List](#)
[Functions](#)
[Reverse](#)

Discuss

[Mailing List](#)
[Twitter](#)
[News](#)

Introduction

Miniboxing is a compilation scheme that improves the performance of generics in [the Scala programming language](#). The miniboxing transformation is generic enough to be potentially useful for any statically typed language running on one of the Java Virtual Machines, such as [Managed X10](#), [Kotlin](#) or [Ceylon](#).

We'll start by following what happens to a generic class, as it gets compiled. Let's take `class C` as an example:

```
class C[T](t: T)
```

After compiling this class to Java Virtual Machine bytecode, under the [erasure transformation](#) one would get bytecode which roughly corresponds to:

```
class C {  
  var t: Object = _      // field  
  def C(t: Object): C = { // constructor  
    this.t = t  
  }  
}
```

As you can see, erasure transformed `t` from a generic value into a pointer to a heap object. While this is perfectly suited for storing a string or another class inside `class C`, it becomes suboptimal when dealing with primitive value types, such as booleans, bytes, integers and floating point numbers.

The reason it's suboptimal is because primitive value types are not heap objects but values passed on the stack, so under the Java Virtual Machine it is common to encode them as heap objects, a process known as [boxing](#).



Start here

- [Start Page](#)
- [Introduction](#)
- [Full Tutorial](#)
- [Benchmarks](#)

Use it

- [Sbt config](#)
- [Command line](#)
- [Source code](#)
- [Issue tracker](#)
- [License](#)

Experiment

- [Sbt project](#)
- [Linked List](#)
- [Functions](#)
- [Reverse](#)

Discuss

- [Mailing List](#)
- [Twitter](#)
- [News](#)

Introduction

Miniboxing is a compilation scheme that improves the performance of generics in the Scala programming language. The miniboxing transformation is generic enough to be potentially useful for any statically typed language running on one of the Java Virtual Machines, such as Managed X10, Kotlin or Ceylon.

We'll start by following what happens to a generic class, as it gets compiled. Let's take `class C` as an example:

```
class C[T](t: T)
```

After compiling this class to Java Virtual Machine bytecode, under the erasure transformation one would get bytecode which roughly corresponds to:

```
class C {  
  var t: Object = _           // field  
  def C(t: Object): C = {     // constructor  
    this.t = t  
  }  
}
```

As you can see, erasure transformed `t` from a generic value into a pointer to a heap object. While this is perfectly suited for storing a string or another class inside `class C`, it becomes suboptimal when dealing with primitive value types, such as booleans, bytes, integers and floating point numbers.

The reason it's suboptimal is because primitive value types are not heap objects but values passed on the stack, so under the Java Virtual Machine it is common to encode them as heap objects, a process known as boxing.

scala-miniboxing.org



Start here

[Start Page](#)
[Introduction](#)
[Full Tutorial](#)
[Benchmarks](#)


Use it

[Sbt config](#)
[Command line](#)
[Source code](#)
[Issue tracker](#)
[License](#)

Experiment

[Sbt project](#)
[Linked List](#)
[Functions](#)
[Reverse](#)

Discuss

[Mailing List](#)
[Twitter](#) 
[News](#) 

Introduction

tutorials

Miniboxing is a transformation scheme that improves the performance of generics in the Scala programming language. The transformation is generic enough to be potentially useful for any statically typed language running on the Java Virtual Machines, such as Managed X10, Kotlin or Ceylon.

We'll start by following what happens to a generic class, as it gets compiled. Let's take `class C` as an example:

```
class C[T](t: T)
```

After compiling this class to Java Virtual Machine bytecode, under the erasure transformation one would get bytecode which roughly corresponds to:

```
class C {  
  var t: Object = _           // field  
  def C(t: Object): C = {     // constructor  
    this.t = t  
  }  
}
```

As you can see, erasure transformed `t` from a generic value into a pointer to a heap object. While this is perfectly suited for storing a string or another class inside `class C`, it becomes suboptimal when dealing with primitive value types, such as booleans, bytes, integers and floating point numbers.

The reason it's suboptimal is because primitive value types are not heap objects but values passed on the stack, so under the Java Virtual Machine it is common to encode them as heap objects, a process known as boxing.

scala-miniboxing.org



Start here

[Start Page](#)
[Introduction](#)
[Full Tutorial](#)
[Benchmarks](#)



Use it

[Sbt config](#)
[Command line](#)
[Source code](#)
[Issue tracker](#)
[License](#)

Experiment

[Sbt project](#)
[Linked List](#)
[Functions](#)
[Reverse](#)

Discuss

[Mailing List](#)
[Twitter](#) 
[News](#) 

Introduction

tutorials

sbt config

Miniboxing is a transformation scheme that improves the performance of generics in the Scala programming language. The transformation is generic enough to be potentially useful for any statically typed language running on top of the Java Virtual Machines, such as Managed X10, Kotlin or Ceylon.

We'll start by following what happens to a generic class, as it gets compiled. Let's take `class C` as an

After compiling this class to Java Virtual Machine bytecode, under the erasure transformation one would get bytecode which roughly corresponds to:

```
class C {  
  var t: Object = _           // field  
  def C(t: Object): C = {     // constructor  
    this.t = t  
  }  
}
```

As you can see, erasure transformed `t` from a generic value into a pointer to a heap object. While this is perfectly suited for storing a string or another class inside `class C`, it becomes suboptimal when dealing with primitive value types, such as booleans, bytes, integers and floating point numbers.

The reason it's suboptimal is because primitive value types are not heap objects but values passed on the stack, so under the Java Virtual Machine it is common to encode them as heap objects, a process known as boxing.

scala-miniboxing.org



Start here

[Start Page](#)
[Introduction](#)
[Full Tutorial](#)
[Benchmarks](#)

Use it

[Sbt config](#)
[Command line](#)
[Source code](#)
[Issue tracker](#)
[License](#)

Experiment

[Sbt project](#)
[Linked List](#)
[Functions](#)
[Reverse](#)

Discuss

[Mailing List](#)
[Twitter](#)
[News](#)

Introduction

tutorials

sbt config

examples

value classes, staging and many more transformations

...on scheme that improves the performance of generics in the Scala programming language. This transformation is generic enough to be potentially useful for any statically typed language, including the Java Virtual Machines, such as Managed X10, Kotlin or Ceylon.

We'll start by following what happens to a generic class, as it gets compiled. Let's take `class C` as an

After compiling this class to Java Virtual Machine bytecode, under the erasure transformation one would get bytecode which roughly corresponds to:

```
// field  
// constructor  
this.t = t  
}
```

As you can see, erasure transformed `t` from a generic value into a pointer to a heap object. While this is perfectly suited for storing a string or another class inside `class C`, it becomes suboptimal when dealing with primitive value types, such as booleans, bytes, integers and floating point numbers.

The reason it's suboptimal is because primitive value types are not heap objects but values passed on the stack, so under the Java Virtual Machine it is common to encode them as heap objects, a process known as boxing.

scala-miniboxing.org



Start here

[Start Page](#)
[Introduction](#)
[Full Tutorial](#)
[Benchmarks](#)


Use it

[Sbt config](#)
[Command line](#)
[Source code](#)
[Issue tracker](#)
[License](#)

Experiment

[Sbt project](#)
[Linked List](#)
[Functions](#)
[Reverse](#)

Discuss

[Mailing List](#)
[Twitter](#) 
[News](#) 

Introduction

tutorials

sbt config

examples

support

value classes, staging and many more transformations

...on scheme that improves the performance of generics in the Scala programming
...transformation is generic enough to be potentially useful for any statically typed
...of the Java Virtual Machines, such as Managed X10, Kotlin or Ceylon.

We'll start by following what happens to a generic class, as it gets compiled. Let's take `class C` as an

After compiling this class to Java Virtual Machine bytecode, under the erasure transformation one would get
bytecode which roughly corresponds to:

```
// field  
= {  
  // constructor  
  this.t = t  
}
```

transformed `t` from a generic value into a pointer to a heap object. While this is
a string or another class inside `class C`, it becomes suboptimal when dealing
primitive value types, such as booleans, bytes, integers and floating point numbers.

The reason it's suboptimal is because primitive value types are not heap objects but values passed on the
stack, so under the Java Virtual Machine it is common to encode them as heap objects, a process known as
boxing.

○ Miniboxing in Practice

○ Documentation

● Quirks

○ Live Demo



Quirks of Scala Specialization

posted by Alex, 03.11.2013.

Having used specialization a lot and having fixed some of its issues, I came across a couple of useful tricks – I want to document them both for myself and others. Specialization is the feature that allows you to generate separate versions of generic classes for primitive types, thus avoiding boxing in most cases. First introduced in Scala 2.8 by Iulian Dragos, by Scala 2.11 specialization has become a pretty robust language feature, and a lot of its issues have been fixed, but there are places where it might stab you in the back if you don't watch out. Problem is, specialization interacts with some edge-cases in the language and obscure language features in ways that are not expected. Sometimes, these are just unresolved bugs. Here are some tips and tricks that might help you.

Note: when used correctly, this is a powerful and extremely useful feature few JVM languages (if any) can parallel these days. Don't get scared by these tips.

Know the conditions for method specialization

Perhaps you're not aware of this, but even if a method is a part of a specialized class and contains specializable code, it will not really be specialized unless the specialized type appears in its argument list or its return type. For example:

```
def getValue[@specialized T]: T = ???

class Foo[@specialized T] {
  var value: T = _
  def reset() {
    value = getValue
  }
}
```

Above, `reset` is not specialized. This is an elaborate design decision taken in specialization. If you want

Quirks of Scala Specialization

posted by Alex, 03.11.2013.

Alex Prokopec, [axel22.github.io](https://github.com/axel22)

Having used specialization a lot and having fixed some of its issues, I came across a couple of useful tricks – I want to document them both for myself and others. Specialization is the feature that allows you to generate separate versions of generic classes for primitive types, thus avoiding boxing in most cases. First introduced in Scala 2.8 by Iulian Dragos, by Scala 2.11 specialization has become a pretty robust language feature, and a lot of its issues have been fixed, but there are places where it might stab you in the back if you don't watch out. Problem is, specialization interacts with some edge-cases in the language and obscure language features in ways that are not expected. Sometimes, these are just unresolved bugs. Here are some tips and tricks that might help you.

Note: when used correctly, this is a powerful and extremely useful feature few JVM languages (if any) can parallel these days. Don't get scared by these tips.

Know the conditions for method specialization

Perhaps you're not aware of this, but even if a method is a part of a specialized class and contains specializable code, it will not really be specialized unless the specialized type appears in its argument list or its return type. For example:

```
def getValue[@specialized T]: T = ???

class Foo[@specialized T] {
  var value: T = _
  def reset() {
    value = getValue
  }
}
```

Above, `reset` is not specialized. This is an elaborate design decision taken in specialization. If you want

Quirks of Scala Specialization

posted by Alex, 03.11.2013.

Alex Prokopec, [axel22.github.io](https://github.com/axel22)

Having used specialization a lot and having fixed some of its issues, I came across a couple of useful tricks – I want to document them both for myself and others. Specialization is the feature that allows you to generate separate versions of generic classes for primitive types, thus avoiding boxing in most cases. First introduced in Scala 2.8 by Iulian Dragos, by Scala 2.11 specialization has become a pretty robust language feature, and a lot of its issues have been fixed, but there are places where it might stab you in the back if you don't watch out. Problem is, specialization interacts with some edge-cases in the language and obscure language features in ways that are not expected. Sometimes, these are just unresolved bugs. Here are some tips and tricks that might help you.

Note: when used correctly, this is a powerful and extremely useful feature few JVM languages (if any) can parallel these days. Don't get scared by these tips.

Know the conditions for method specialization

Perhaps you're not aware of this, but even if a method is a part of a specialized class and contains specializable code, it will not really be specialized unless the specialized type appears in its argument list or its return type. For example:

```
def getValue[@specialized T]: T = ???

class Foo[@specialized T] {
  var value: T = _
  def reset() {
    value = getValue
  }
}
```

Above, `reset` is not specialized. This is an elaborate design decision taken by specialization. If you want

Quirks of Scala Specialization

posted by Alex, 03.11.2013.

Alex Prokopec, [axel22.github.io](https://github.com/axel22)

Having used specialization a lot and having fixed some of its issues, I came across a couple of useful tricks – I want to document them both for myself and others. Specialization is the feature that allows you to generate separate versions of generic classes for primitive types, thus avoiding boxing in most cases. First introduced in Scala 2.8 by Iulian Dragos, by Scala 2.11 specialization has become a pretty robust language feature, and a lot of its issues have been fixed, but there are still places where it might catch you in the back if you don't watch out. Problem is, specialization interacts with a some edge-cases in the language and obscure language features in ways that are not expected. Sometimes, these are just unsolvable bugs, here are some tips and tricks that might help you.

Know the conditions for method specialization - method specialization

Know the conditions for method specialization - class specialization

Initialize specialized values outside constructor body

Note: which does correctly, this is a powerful and extremely useful feature. Now JVM languages (if any) can parallel these days. Don't get scared by these tips.

Resolve access problems using the package-private modifier

Use traits where possible

Avoid traits where possible

Perhaps you're not aware of this, but even if a method is a part of a specialized class and contains specialized code, the specialized type appears in its argument list or its return type. For example:

Make your classes as flat as possible

Avoid super calls

Be wary of vars

Think about the primitive types you really care about

Avoid using specialization and implicit classes

```
def getValue[@specialized T]: T = ???  
  
class Foo[@specialized T] {  
  var value: T = _  
  def reset() {  
    value = getValue  
  }  
}
```

Above, `reset` is not specialized. This is an elaborate design decision taken for specialization. If you want

Quirks of Scala Specialization

posted by Alex, 03.11.2013.

Alex Prokopec, [axel22.github.io](https://github.com/axel22)

Having used specialization a lot and having fixed some of its issues, I came across a couple of useful tricks – I want to document them both for myself and others. Specialization is the feature that allows you to generate separate versions of generic classes for primitive types, thus avoiding boxing in most cases. First introduced in Scala 2.8 by Iulian Dragos, by Scala 2.11 specialization has become a pretty robust language feature, and a lot of its issues have been fixed, but there are still places where it might catch you in the back if you don't watch out. Problem is, specialization interacts with a some edge-cases in the language and obscure language features in ways that are not expected. Sometimes, these are just unresolvable bugs, here are some tips and tricks that might help you.

Know the conditions for method specialization - method specialization

Know the conditions for method specialization - class specialization

Initialize specialized values outside constructor body

Resolve access problems using the package-private modifier

Use traits where possible

Avoid traits where possible

Make your classes as flat as possible

Avoid super calls

Be wary of vars

Think about the primitive types you really care about

Avoid using specialization and implicit classes

```
def getValue[@specialized T]: T = ???
```

```
class Foo[@specialized T] {
```

```
  var value: T = _
```

```
  def reset() {
```

```
    value = getValue
```

```
  }
```

```
}
```



Quirks of Scala Specialization

posted by Alex, 03.11.2013.

Alex Prokopec, axel22.github.io


Having used specialization a lot and having fixed some of its issues, I came across a couple of useful tricks – I want to document them both for myself and others. Specialization is the feature that allows you to generate separate versions of generic classes for primitive types, thus avoiding boxing in most cases. First introduced in Scala 2.8 by Iulian Dragos, by Scala 2.11 specialization has become a pretty robust language feature, and a lot of issues have been fixed, but there are still places where it might catch you in the back if you don't watch out. Problem is, specialization interacts with a some edge-cases in the language and obscure language features in ways that are not expected. Sometimes, these are just unsolved things, here are some tips and tricks that might help you.

Now, which one is correct? This is a powerful and extremely useful feature. New JVM languages (if any) can parallel these days. Don't get scared by these tips.

Know the conditions for method specialization

Perhaps you're not aware of this, but even if a method is a part of a specialized class and contains specializations, the specialized type appears in its argument list or its return type. For example:

```
def getValue[@specialized T]: T = ???  
  
class Foo[@specialized T] {  
  var value: T = _  
  def reset(): T = {  
    value = getValue  
  }  
}
```



Quirks of Scala Specialization

posted by Alex, 03.11.2013.

Alex Prokopec, axel22.github.io

Having used specialization a lot and having fixed some of its issues, I came across a couple of useful tricks – I want to document them both for myself and others. Specialization is the feature that allows you to generate separate versions of generic classes for primitive types, thus avoiding boxing in most cases. It was first introduced in Scala 2.8 by Iulian Dragos, by Scala 2.11 specialization has become a pretty robust language feature, and a lot of issues have been fixed, but there are still places where it might catch you in the back if you don't watch out. Problem is, specialization interacts with a some edge-cases in the language and obscure language features in ways that are not expected. Sometimes these are just unsolved bugs, here are some tips and tricks that might help you.

● = fixed by the 0.4 release

● Know the conditions for method specialization - method specialization

● Know the conditions for method specialization - class specialization

● Initialize specialized values outside constructor body

● Resolve access problems using the package-private modifier

● Use traits where possible

● Avoid traits where possible

● Make your classes as flat as possible

● Avoid super calls

● Be wary of vars

● Think about the primitive types you really care about

● Avoid using specialization and implicit classes

```
def getValue[@specialized T]: T = ???  
  
class Foo[@specialized T] {  
  var value: T = _  
  def reset() {  
    value = getValue  
  }  
}
```



Quirks of Scala Specialization

posted by Alex, 03.11.2013.

Alex Prokopec, axel22.github.io

Having used specialization a lot and having fixed some of its issues, I came across a couple of useful tricks – I want to document them both for myself and others. Specialization is the feature that allows you to generate separate versions of generic classes for primitive types, thus avoiding boxing in most cases. First introduced in Scala 2.8 by Iulian Dragos, by Scala 2.11 specialization has become a pretty robust language feature, and a lot of its issues have been fixed, but there are places where it might stab you in the back if you don't watch out. Problem is, specialization interacts with some edge-cases in the language and obscure language features in ways that are not expected. Sometimes, these are just unresolved bugs. Here are some tips and tricks that might help you.

Quirk

Note: when used correctly, this is a powerful and extremely useful feature few JVM languages (if any) can parallel these days. Don't get scared by these tips.

Know the conditions for method specialization

Perhaps you're not aware of this, but even if a method is a part of a specialized class and contains specializable code, it will not really be specialized unless the specialized type appears in its argument list or its return type. For example:

```
def getValue[@specialized T]: T = ???

class Foo[@specialized T] {
  var value: T = _
  def reset() {
    value = getValue
  }
}
```



Above, `reset` is not specialized. This is an elaborate design decision taken by specialization. If you want

Quirks of Scala Specialization

posted by Alex, 03.11.2013.

Alex Prokopec, [axel22.github.io](https://github.com/axel22)

Having used specialization a lot and having fixed some of its issues, I came across a couple of useful tricks – I want to document them both for myself and others. Specialization is the feature that allows you to generate separate versions of generic classes for primitive types, thus avoiding boxing in most cases. First introduced in Scala 2.8 by Iulian Dragos, by Scala 2.11 specialization has become a pretty robust language feature, and a lot of its issues have been fixed, but there are places where it might stab you in the back if you don't watch out. Problem is, specialization interacts with some edge-cases in the language and obscure language features in ways that are not expected. Sometimes, these are just unresolved bugs. Here are some tips and tricks that might help you.

Quirk = limitation

Note: when used correctly, this is a powerful and extremely useful feature few JVM languages (if any) can parallel these days. Don't get scared by these tips.

Know the conditions for method specialization

Perhaps you're not aware of this, but even if a method is a part of a specialized class and contains specializable code, it will not really be specialized unless the specialized type appears in its argument list or its return type. For example:

```
def getValue[@specialized T]: T = ???

class Foo[@specialized T] {
  var value: T = _
  def reset() {
    value = getValue
  }
}
```



Above, `reset` is not specialized. This is an elaborate design decision taken by specialization. If you want

Quirks of Scala Specialization

posted by Alex, 03.11.2013.

Alex Prokopec, axel22.github.io

Having used specialization a lot and having fixed some of its issues, I came across a couple of useful tricks – I want to document them both for myself and others. Specialization is the feature that allows you to generate separate versions of generic classes for primitive types, thus avoiding boxing in most cases. First introduced in Scala 2.8 by Iulian Dragos, by Scala 2.11 specialization has become a pretty robust language feature, and a lot of its issues have been fixed, but there are places where it might stab you in the back if you don't watch out. Problem is, specialization interacts with some edge-cases in the language and obscure language features in ways that are not expected. Sometimes, these are just unresolved bugs. Here are some tips and tricks that might help you.

Quirk = limitation + silent failure

Note: when used correctly, this is a powerful and extremely useful feature few JVM languages (if any) can parallel these days. Don't get scared by these tips.

Know the conditions for method specialization

Perhaps you're not aware of this, but even if a method is a part of a specialized class and contains specializable code, it will not really be specialized unless the specialized type appears in its argument list or its return type. For example:

```
def getValue[@specialized T]: T = ???

class Foo[@specialized T] {
  var value: T = _
  def reset() {
    value = getValue
  }
}
```



Above, `reset` is not specialized. This is an elaborate design decision taken by specialization. If you want

Quirks of Scala Specialization

posted by Alex, 03.11.2013.

Alex Prokopec, axel22.github.io

Having used specialization a lot and having fixed some of its issues, I came across a couple of useful tricks – I want to document them both for myself and others. Specialization is the feature that allows you to generate separate versions of generic classes for primitive types, thus avoiding boxing in most cases. First introduced in Scala 2.8 by Iulian Dragos, by Scala 2.11 specialization has become a pretty robust language feature, and a lot of its issues have been fixed, but there are places where it might stab you in the back if you don't watch out. Problem is, specialization interacts with some edge-cases in the language and obscure language features in ways that are not expected. Sometimes, these are just unresolved bugs. Here are some tips and tricks that might help you.

Quirk = limitation + silent failure

Note: when used correctly, this is a powerful and extremely useful feature few other languages (if any) can parallel these days. Don't get scared by these tips.

let's make the limitations transparent

Know the conditions for method specialization

Perhaps you're not aware of this, but even if a method is a part of a specialized class and contains specializable code, it will not really be specialized unless the specialized type appears in its argument list or its return type. For example:

```
def getValue[@specialized T]: T = ???

class Foo[@specialized T] {
  var value: T = _
  def reset() {
    value = getValue
  }
}
```




Above, `reset` is not specialized. This is an elaborate design decision taken for specialization. If you want

○ Miniboxing in Practice

○ Documentation

○ Quirks

● Live Demo






Miniboxing

Theory

Practice

Benchmarks

Conclusion



scala-miniboxing.org

○ Benchmarks

○ Microbenchmarks


○ Collections

○ Spire




Miniboxing Benchmarks


on array buffers



Miniboxing Benchmarks on linked lists




Miniboxing Benchmarks on linked lists



Miniboxing Benchmarks

on the Spire library (RexBench)





Miniboxing

Theory

Practice

Benchmarks

Conclusion





Credits and Thank you-s

- Cristian Talau - developed the initial prototype, as a semester project
- Eugene Burmako - the value class plugin based on the LDL transformation
- Aymeric Genet - developing collection-like benchmarks for the miniboxing plugin
- Martin Odersky, for his patient guidance
- Eugene Burmako, for trusting the idea enough to develop the value-plugin based on the LDL transformation
- Iulian Dragos, for his work on specialization and many explanations
- Miguel Garcia, for his original insights that spawned the miniboxing idea
- Michel Schinz, for his wonderful comments and enlightening ACC course
- Andrew Myers and Roland Ducournau for the discussions we had and the feedback provided
- Heather Miller for the eye-opening discussions we had
- Vojin Jovanovic, Sandro Stucki, Manohar Jonalagedda and the whole LAMP laboratory in EPFL for the extraordinary atmosphere
- Adriaan Moors, for the miniboxing name which stuck :))
- Thierry Coppey, Vera Salvisberg and George Nithin, who patiently listened to many presentations and provided valuable feedback
- Grzegorz Kossakowski, for the many brainstorming sessions on specialization
- Erik Osheim, Tom Switzer and Rex Kerr for their guidance on the Scala community side
- OOPSLA paper and artifact reviewers, who reshaped the paper with their feedback
- Sandro, Vojin, Nada, Heather, Manohar - reviews and discussions on the LDL paper
- Hubert Plociniczak for the type notation in the LDL paper
- Denys Shabalin, Dmitry Petrashko for their patient reviews of the LDL paper

Special thanks to the Scala Community for their support!

(@StuHood, @vpatryshev and everyone else!)






scala-miniboxing.org



ScalaTeam @ EPFL



scala-miniboxing.org

ScalaTeam @ EPFL

- **YinYang** frontend multi-stage execution
 - based on macro transformations
 - Vojin Jovanovic/Sandro Stucki



<https://github.com/vjovanov/yin-yang>

ScalaTeam @ EPFL



- **Scala.js** backend
 - compiles Scala to JavaScript
 - Sébastien Doeraene/Tobias Schlatter



<http://www.scala-js.org/>

ScalaTeam @ EPFL

- **Lightweight Modular Staging**
 - program optimization
 - Tiark Rompf + pretty much everyone



<http://scala-lms.github.io/>

ScalaTeam @ EPFL

- **Dependent Object Types** calculus
 - core type system of the **dotty** compiler
 - Nada Amin/Tiark Rompf



<https://github.com/TiarkRompf/minidot>



<https://github.com/lampepfl/dotty>

ScalaTeam @ EPFL

- **Pickling** framework and **Spores**
 - support for **distributed programming**
 - Heather Miller/Philipp Haller + others



<https://github.com/scala/pickling>



<https://github.com/heathermiller/spores>

ScalaTeam @ EPFL

- **Staged Parser-combinators**
 - fast parser combinators through staging
 - Manohar Jonnalagedda + others



<https://github.com/manojo/experiments>

ScalaTeam @ EPFL

- **dotty** compiler
 - compiler for Scala but with the DOT type system
 - Martin Odersky/Dmitry Petrashko/Tobias Schlatter



<https://github.com/lampepfl/dotty>

ScalaTeam @ EPFL

- **scala.meta** metaprogramming support
 - Improved reflection, macros, and many more
 - Eugene Burmako/Denys Shabalin + others



<http://scalameta.org/>

ScalaTeam @ EPFL



- **scaladyno** plugin
 - giving Scala a dynamic language look and feel
 - Cédric Bastin/Vlad Ureche



<https://github.com/scaladyno/scaladyno-plugin>

ScalaTeam @ EPFL



- **miniboxing** specialization
 - LDL transformation
 - Vlad Ureche/Aymeric Genêt + others

[www http://scala-miniboxing.org/](http://scala-miniboxing.org/)

ScalaTeam @ EPFL



- **ScalaBlitz** optimization framework
 - macro-based collection optimization
 - Dmitry Petrashko/Aleksandar Prokopec

[www http://scala-blitz.github.io/](http://scala-blitz.github.io/)

ScalaTeam @ EPFL

- **LMS-Kappa** protein simulator
 - using multi-stage programming for performance
 - Sandro Stucki



<https://github.com/sstucki/lms-kappa>

ScalaTeam @ EPFL

- **Odds** probabilistic programming framework
 - using scala-virtualized and macros
 - Sandro Stucki



<https://github.com/sstucki/odds>

ScalaTeam @ EPFL

- **Type debugger** for Scala
 - debugging aid for Scala type errors
 - Hubert Plociniczak



[http://lampwww.epfl.ch/~plocinic/
type-debugger-tutorial/](http://lampwww.epfl.ch/~plocinic/type-debugger-tutorial/)

ScalaTeam @ EPFL



- **ScalaMeter** benchmarking framework
 - google caliper for scala
 - Aleksandar Prokopec



<http://scalameter.github.io/>

ScalaTeam @ EPFL

- **Vector** implementation using RRB trees
 - improved performance for Scala collections
 - Nicolas Stucki



<https://github.com/nicolasstucki/scala-rrb-vector>

ScalaTeam @ EPFL

- the **new scalac backend**
 - good performance gains
 - Miguel Garcia
 - on the job market right now



<http://magarciaepfl.github.io/scala/>



miguel.garcia@tuhh.de